


Introduction to MongoDB



Why NoSql ?

- Relational databases are not designed to scale
- schema, joins

But people want:

- Scale
- Speed
- Cloud
- New data

C and Latency Tradeoff

- Amazon claims that just an extra one tenth of a second on their response times will cost them 1% in sales.
- Google said they noticed that just a half a second increase in latency caused traffic to drop by a fifth.

What is NoSQL?

- non-relational
- simple API
- schema-free
- open-source
- horizontally scalable (sharding)
- replication support
- eventually consistent /BASE

Different types of NoSQL Databases

- NoSQL database are classified according to their data storage models:
 - Column (Cassandra)
 - Document (MongoDB)
 - Key – value Pair(Dynamo – Amazon)
 - Graph

MongoDB

- Name derived from Hu(**MONGO**)us word
- Document Oriented Database
- Built for High – Performance and scalability
- Document based queries for **Easy Readability**
- Replication and failover for **High Availability**
- Auto Sharding for **Easy Scalability**

Comparison between RDBMS and NoSQL DB

- Example: Class
- Location
- Presenter
 - Presenting at a location
- People
 - Potential attendees in context of a class
- Class
 - Presenter in location with people as actual attendees

Relational Database: Example

- Class schema in a relational database
- Presentation { id, name, location }
- People { id, name }
- Address { id, city, state, zip }

Schema for this class in a relational database model

Presentation			Address		
id name location			id city state		
1	Chris	SJSU	SJSU	San Jose	CA

People		Class	
id name		id person presentation	
10	Simon	20	10 1
11	Chris	20	11 1

Relational database: Example

```
CREATE TABLE Presentation (  
    id Integer primary key, name String, location string,  
    FOREIGN KEY (location) REFERENCES Address(id));  
CREATE TABLE Address (  
    id String primary key, city String, state String);  
CREATE TABLE People (  
    id Integer primary key, name String);  
CREATE TABLE Class (  
    id Integer, person Integer, presentation Integer,  
    PRIMARY KEY (id, person, presentation),  
    FOREIGN KEY (person) REFERENCES People(id),  
    FOREIGN KEY (presentation) REFERENCES Presentation(id));
```

Relational database: Example

```
select Presentation.name, Presentation.location,  
       Address.city, Address.state, People.name  
from Presentation, Address, People, Class  
where Class.person = People.id  
       and Class.presentation = Presentation.id  
       and Presentation.location = Address.id;
```

name	location	city	state	name
Chris	SJSU	San Jose	CA	Simon
Chris	SJSU	San Jose	CA	Chris

Relational Database: Recap

1. Schema design

Primary key (underlined) and foreign key (cursive) constraints

2. Table creation

DDL

3. Data insertion for each table

DML

4. Query: join

DML

5. Data structure creation within application system

JDBC resultset to e.g. Java objects

NoSQL Database: Use Case Example

```
use course /* database will be created if not present */
db.presentation.insert(
  {"id": 1,
    "name": "Simon",
    "location": {"id": "SJSU",
                  "city": "San Jose",
                  "state": "CA"},
    },
  "people": [{"id": 10, "name": "Simon"},
              {"id": 11, "name": "Chris"}
            ]
  )
```

NoSQL Database: Use Case Example

- `db.presentation.find()`
- `db.presentation.find({ "id": 1 })`

NoSQL Database: Recap

1. Schema design

Primary key (underlined) and foreign key (cursive) constraints

2. Table creation

DDL

3. Data insertion for each table

DML

4. Query: join

DML



5. Data structure creation within application system

IDBC resultset to e.g. Java objects

NoSQL Database: Major Players

- Too many document NoSQL databases to name a few distinct ones

29 systems in ranking, July 2014

Rank	Last Month	DBMS	Database Model	Score	Changes
1.	1.	MongoDB	Document store	238.78	+7.33
2.	2.	CouchDB	Document store	23.07	+0.28
3.	3.	Couchbase	Document store	16.58	+0.79
4.	4.	MarkLogic	Multi-model 	8.20	-0.02
5.	5.	RavenDB	Document store	5.09	-0.42
6.	6.	GemFire	Document store	2.16	-0.06
7.	7.	OrientDB	Multi-model 	1.71	-0.02
8.	8.	Cloudant	Document store	1.70	+0.07
9.	9.	Datameer	Document store	0.88	+0.08
10.	10.	Mnesia	Document store	0.72	+0.01

Key Benefit of NoSQL: $O(1)$ Lookup

- Fast lookup
 - No joining required
 - All data about one domain concept in one document
- Direct programming language representation
 - No mapping or ‘ORM’ layer required
- JSON library
 - Direct result representation and manipulation
 - JavaScript: representation in language data types directly
 - E.g., check out MongoDB node.js driver

Key Problem of NoSQL: No Join Operator

- Many NoSQL databases do not implement a join query operator
 - If you need to join data, then you have to do it in the application system layer
- But, wait a moment ...
 - Is it ever necessary to join data in NoSQL databases?
 - Some claim: not necessary due to support of
 - Sub-documents
 - Arrays (lists)
- Let's look at an example
 - Supplier - Parts

Key Problem of NoSQL: No Join Operator

- Example
 - Supplier - Parts relationship (N:M)
 - Each supplier supplies many parts
 - Each part supplied by many suppliers
- Relational DBMS
 - “Supplier” table
 - “Part” table
 - “Supplies” relationship in table

Key Problem of NoSQL: No Join Operator

Supplier - Part - Supplies

Supplier		Part		Supplies	
id	name	id	name	<i>supplier_id</i>	<i>part_id</i>
10	Supp1	20	Part1	10	20
11	Supp2	21	Part2	10	21
				11	20

Key Problem of NoSQL: No Join Operator

Supplier - Supplies – Part

```
{ "id": 10,  
  "name": "Supp1",  
  "supplies": [{ "id": 20, "name": "Part1"},  
                { "id": 21, "name":  
                  "Part2"} ] }  
  
{ "id": 11,  
  "name": "Supp2",  
  "supplies": [{ "id": 20, "name": "Part1"} ] }
```

Supplier - Supplies – Part

```
{ "id": 10,  
  "name": "Supp1",  
  "supplies": [20, 21] }  
  
{ "id": 10,  
  "name": "Supp1",  
  "supplies": [20, 21] }  
  
{ "id": 20, "name": "Part1" }  
{ "id": 21, "name": "Part2" }
```

Why use MongoDB?

- MongoDB stores data in Objects
- Uses BSON (Binary JSON)
- No Joins
- No Complex Queries
- Embedded Documents and arrays reduce the need for joins
- No multi-document transactions

Where to use MongoDB ?

- Ideal for Web Applications
- Applications containing semi-structured data and need flexible schema management
- Caching and High Scalability
- Scenarios where **data availability** and **size of data** are priorities over the **transactions** of data

When to not use MongoDB?

- ACID properties are important for storage
- Highly Transactional Applications (Banking domain, Security)
- Problems and applications requiring Joins and complex queries

Key Problem of NoSQL: No Database-Enforced Consistency

- Not enforced
 - Primary key
 - Foreign key
 - Enumeration
 - Cascading delete
 - etc.
- Enforcement can be accomplished
 - When
 - reading or writing
 - In application system code
 - In self-implemented database access layer
 - In separate consistency check process
 - Not at all

How does MongoDB Store data?

- Stores data in form of Documents
- JSON like field – value pair
- Documents analogous to structures in programming languages with key – value pair
- Documents stored in **BSON (Binary JSON)** format
- BSON is JSON with additional type information

NoSQL: Key Insights

- Specialized data models
 - Not universal, but optimized towards special cases
- Specialized query access
 - Not universal, but optimized towards special cases
- Different / absent consistency supervision
 - Relaxed constraints
- Trade-off
 - Gain through specialization
 - Implementation of missing functionality outside of database


CAP Theorem: Theory

The **CAP theorem** states that it is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- *Consistency*
- *Availability*
- *Partition tolerance*

CAP Theorem

- Consistency
- Availability
- Partition Tolerance
- Choose two



Questions?

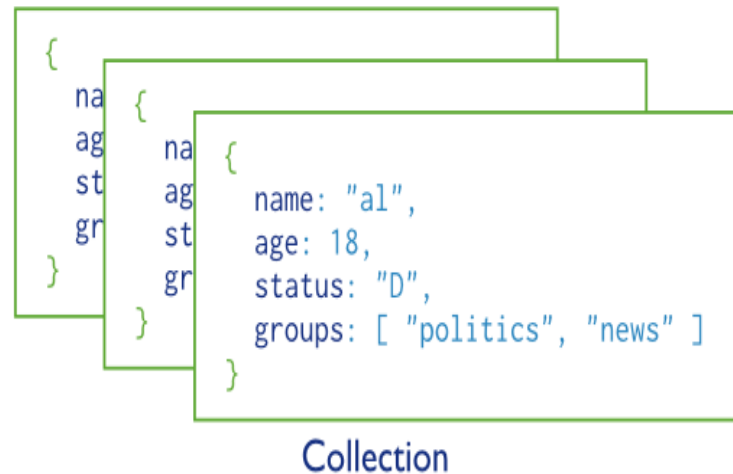
- Which one would you choose when network partition?
(a) C (b) A
- Which of CAP is essential for a distributed system?
(a) C (b) A (c) P (d) none of the above
- What is missing in The CAP Theorem in implementing distributed systems?

CAP

- Dynamo does not guarantee C by default
- The event of P forces systems to decide on reducing C or A
- What is the probability of P?
 - Local network
 - Wide area network

Collections in MongoDB

- MongoDB stores all data in Collections
- Collections in MongoDB analogous to tables in relational databases
- It is schema – less and contains a group of related documents
- Created on-the-fly when referenced for the first time



Document in MongoDB

- Stored in Collections
- Analogous to Records/Rows in Relational databases
- Has **_id** field – works like Primary keys in Relational databases
- Sample document containing name, age, status and groups

groups{

```
name: "sue",  
age: 26,  
status: "A",  
groups: [ "news", "sports" ]
```

← field: value
← field: value
← field: value
← field: value

}

Queries in MongoDB

- MongoDB provides **db.collection.find()** method to retrieve data
- This method accepts both query criteria and

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
).limit(5)
```

← collection
← query criteria
← projection
← cursor modifier

Mongo Query

```
SELECT _id, name, address  
FROM users  
WHERE age > 18  
LIMIT 5
```

Similar SQL Query

← table
← select criteria
← cursor modifier

Projections - Queries in MongoDB

- If you include 1 – in projection parameter, it returns the value
- If you include 0 – in projection parameter, it eliminates it from the result

```
db.records.find( { "user_id": { $lt: 42 } }, { "_id": 0, "name": 1 , "email": 1 } )
```

- `_id` – always included in results. Specify “`_id : 0`” to exclude it from results

```
db.records.find( { "user_id": { $lt: 42 } }, { "history": 0 } )
```

- Excludes history field from the results, and returns all other fields

Insert Operation

- In MongoDB, `db.collection.insert()` method adds new documents to collections

```
db.users.insert (  ← collection
{
  name: "sue",      ← field: value
  age: 26,          ← field: value
  status: "A"       ← field: value
}                  } document
)
```

Mongo Insert

SQL insert

```
INSERT INTO users      ← table
  ( name, age, status ) ← columns
VALUES ( "sue", 26, "A" ) ← values/row
```

Insert Operation

- If `_id` is not included in the insert query, mongo adds `_id` internally and computes its value with a unique **ObjectId**
- **ObjectId:**
 - 12 byte BSON type
 - Guarantees uniqueness within that collection
 - Generated based on timestamp, machine ID, process ID and a internal process-local incremental counter

Update Operation

- In MongoDB, `db.collection.update()` method modifies existing documents in a collection

```
db.users.update(  
  { age: { $gt: 18 } },  
  { $set: { status: "A" } },  
  { multi: true }  
)
```

← collection
← update criteria
← update action
← update option

Mongo Update

SQL update

```
UPDATE users  
SET    status = 'A'  
WHERE  age > 18
```

← table
← update action
← update criteria

Update Operation

```
db.users.update(  
  { age: { $gt: 18 } },  
  { $set: { status: "A" } },  
  { multi: true }  
)
```

← collection
← update criteria
← update action
← update option

- Updates on **users** collection
- Sets “**status**” field to “A”
- With criteria of “**age**” greater than “**18**”
- **multi: true** – updates all the document in a query with the matching criteria

Remove Operation

- In MongoDB, `db.collection.remove()` method deletes document from the collection

```
db.users.remove(  ← collection  
  { status: "D" } ← remove criteria  
)
```

Mongo Delete

SQL Delete

```
DELETE FROM users ← table  
WHERE status = 'D' ← delete criteria
```


Remove Operation

```
db.users.remove(    ← collection  
  { status: "D" }    ← remove criteria  
)
```

- Delete operation performed on users collection
- Removes all documents with “**status**” as “**D**”

Additional Operations

- **db.collections.save()**
 - Updates an existing documents if it finds the document with the mentioned values
 - Inserts in the collection, if it does not find a document with the mentioned values

Installing MongoDB

- In Windows:

<https://www.youtube.com/watch?t=1&v=sBdaRlgb4N8>

- In Mac/Linux:

https://www.youtube.com/watch?v=_WJ8m5QHvwc

Using MongoDB with Node.js


- **Install MongoDB Node.js Module**

```
npm install mongodb
```

Example

- Login Application
- Access MongoDB to authentic the user
- Use Mongo Store to store sessions in
MongoDB

Example – Login Page

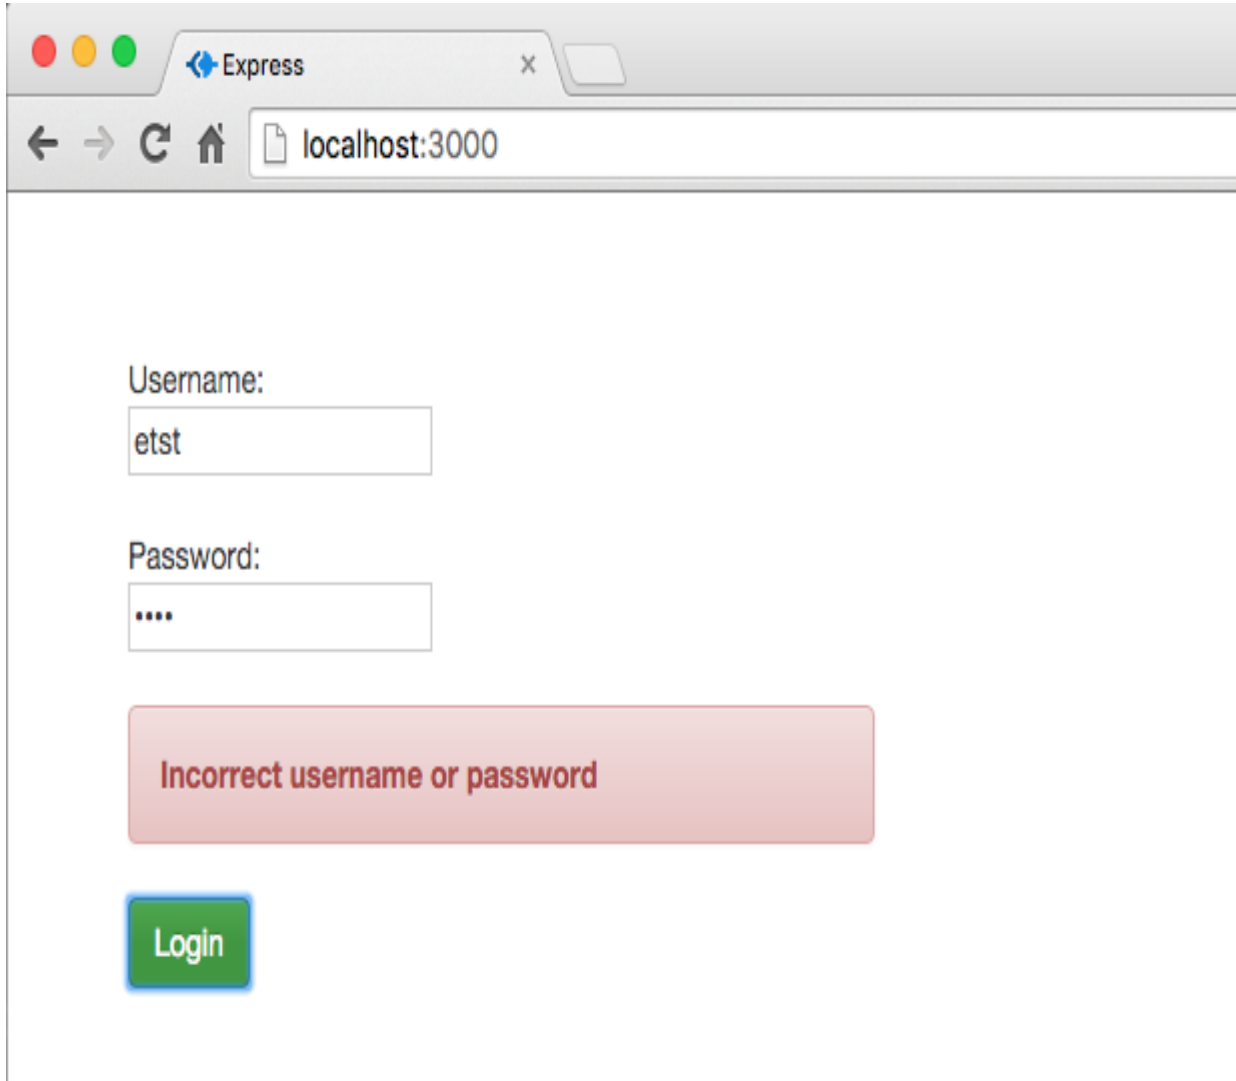


Username:

Password:

Login

Example – Wrong Credentials



A screenshot of a web browser window. The browser's address bar shows 'localhost:3000'. The page contains a login form with two input fields: 'Username:' with the value 'etst' and 'Password:' with masked characters '....'. Below the password field is a red error message box that says 'Incorrect username or password'. At the bottom of the form is a green 'Login' button.


Username:
etst

Password:
....

Incorrect username or password

Login

Example - Homepage



Welcome to the Portal, test

Logout

Exercise

- Create a Login Application
- Should have option to sign up the user
- Login with the same user
- Show the details of the signed in user
- Use MongoDB to store the data
- Use Passport with Express session for authentication

References

- SQL vs NoSQL - <https://www.mongodb.com/nosql-explained>
- MongoDB Introduction - <http://docs.mongodb.org/manual/core/crud-introduction/>
- Installing MongoDB (Mac) - https://www.youtube.com/watch?v=_WJ8m5QHvwc
- Installing MongoDB (Windows) - <https://www.youtube.com/watch?t=1&v=sBdaRlgb4N8>